


# Kafka Workshop [Wefair]

12.3.19

- Logs
- Btree
- Pub/Sub


- Sequential reading is distributed
- Should not have 2 consumers / partitions.



Kafka


- Horizontal Scalability

→ Consistent Hashing




- Disk partition
- Global partition (row-topic wise)

- Consumers can read @ own speed.



- Messages will disappear after retention.





Replication

- Log - Type of durable messaging system
  - like Rabbit , ActiveMQ
  - Scalable
  - Built in fault tolerance
  - Storage
- { Jabscon  
Test }

Demo


- Connectors
- UI to monitor Kafka

Stream Processing :



eg

Create Stream possible as  
Select card-no, count(\*) from auth-attempt  
window



Producer Consumer ↓ Flexibility

Viny Njuma

1. Track - Clickstream

Frogem  
Siphon

• Clickstream → Tracking on site


⇒ Recommendation System.

• Loggy 8 metrics →

Populating ELK 8 influx workstreams




• Orwell

Local Aggregate model



• minotmaker





### Monitoring Kafka

- KPI
- URP
- Offline Partition
- Disk Usage
- msg Rates

### Monitoring KPI

- mm Process
- Bytes Out
- Bytes In
- Mem Usage
- Load Avg

### Kafka Manager

Administering Kafka Admin

- Newm... v
- Kafka Minimizer
- Replace RabbitMQ Workflow.
- Custom Scoring

v narayan@wafair

Last modified: 7:44 PM